

Gesture UI Machine Learning Tensorflow

Donny Hurley

Atlantic Technological University



TensorFlow is an open-source symbolic maths library that is popular in Machine Learning, especially for Neural Networks

Developed by GoogleBrain, Google's AI/Deep learning division

Note that only the implementation of TensorFlow that we see on GitHub is open-source. Google maintains its own internal version.

It is not the first package. What is unique about TensorFlow is how quickly it has gained so much market share. It is the dominant framework for NNs

- Theano
- Caffe
- Torch
- and there are others

What they all have in common is they are designed to complete fast mathematics problems, in particular linear algebra routines for arrays (multiplication of matrices in particular), which is the backbone of Machine Learning techniques

Keras is an open-source deep learning library written in Python.

Developing deep learning models using mathematical libraries like TensorFlow, Theano, and PyTorch was cumbersome, requiring tens or even hundreds of lines of code to achieve the simplest tasks. The focus of these libraries was on research, flexibility, and speed, not ease of use.

Note

Keras is an API designed for human beings, not machines. Keras follows best practices for reducing cognitive load: it offers consistent & simple APIs, it minimizes the number of user actions required for common use cases, and it provides clear & actionable error messages. It also has extensive documentation and developer guides.

The API is clean and simple, allowing standard deep learning models to be defined, fit, and evaluated in just a few lines of code.

A secondary reason Keras took-off was because it allowed you to use any one among the range of popular deep learning mathematical libraries as the backend (e.g. used to perform the computation), such as TensorFlow, Theano, and later, CNTK. This allowed the power of these libraries to be harnessed (e.g. GPUs) with a very clean and simple interface.

In 2019 TensorFlow 2 was released by Google that made a lot of changes to the API, some fundamental changes and the code is mostly not compatible

Be careful when looking up tutorials, which version of TensorFlow are they made for.

<https://www.datasciencecentral.com/profiles/blogs/tensorflow-1-x-vs-2-x-summary-of-changes> shows a summary of changes.

In short

- TensorFlow 2 tries to be more "Pythonic" and easy to use
- By default it uses Eager execution (to be more consistent with how Python usually works), whereas 1 used Sessions to run.
- Has built in support for Keras
 - This is usually referred to as `tf.keras`

For building models so much less complex

Commonly the Keras API integrated with TensorFlow is referred to as `tf.keras`. This is to distinguish it from the standalone Keras open source project that supports other backends and it follows the usual way of importing

```
import tensorflow as tf
tf.keras....
```

Quote

At this time, we recommend that Keras users who use multi-backend Keras with the TensorFlow backend switch to `tf.keras` in TensorFlow 2.0. `tf.keras` is better maintained and has better integration with TensorFlow features (eager execution, distribution support and other).

from Keras homepage.

See [PyImageSearch](#)

for a comparison and pay note to

Quote

Keras v2.3.0 is the first release of Keras that brings keras in sync with tf.keras. It will be the the last major release to support backends other than TensorFlow (i.e., Theano, CNTK, etc.)

And most importantly, deep learning practitioners should start moving to TensorFlow 2.0 and the tf.keras package

Install TensorFlow

It does not come as default through Anaconda so you will need to install it using conda/pip or whatever method you use.

Your scripts should always verify the installation and version number

```
import tensorflow as tf
print(tf.__version__)
```

to ensure it is installed correctly and the version is what you expect.

Often you will get warning messages when using the API about what your hardware supports but it was not configured for use. This could be things like CPU instructions/GPU instructions.

They are for information and will not prevent the code from running. However, some instructions can speed up the Maths that the library does. GPUs are especially efficient at Matrix maths.

Install TensorFlow

<https://www.tensorflow.org/install/pip>

has instructions for installing it through pip

<https://docs.anaconda.com/anaconda/user-guide/tasks/tensorflow/> has

instructions for installing it through conda

If you want it GPU accelerated you will need CUDA if you've an Nvidia GPU and instructions are given there too.

```
import tensorflow as tf
if tf.test.gpu_device_name():
    print('Default GPU Device:
          {}'.format(tf.test.gpu_device_name()))

else:
    print("Please install GPU version of TF")
```

A great way to try TensorFlow (and other Python projects) is by using Google collab. It allows you to access GPU resources which can greatly speed up your calculations

Google Colab

Free to use, you will be in a queue for some things, but if you don't have a decent GPU in your own machine can be very useful

Longer Example

Let's use LinearRegression as a basic example to look at building a model and to demonstrate how to use TensorFlow

In short, say we have a sample of data, each item in the sample has 1 feature (x) and you wish to come up with a model

$$\hat{y} = \alpha + \beta x$$

or if you have multiple features (x_1, x_2, \dots, x_m) then the model

$$\hat{y} = \alpha + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_m x_m$$

The process for regression usually follows the following steps

- ① Formulating the problem
 - Selecting a tentative form of the model (linear, polynomial etc)
 - Selecting the features
- ② Fit the model
 - Use the data we know the correct answer to and find the best fit for the parameters with that data.
 - So that the predictions at the end are close to what the actual values were for each x .
- ③ Model validation
 - Evaluate how it performs and the stability of it

Previous slide is fine for any ML process. "Selecting tentative form" could be select a neural network. In terms of how this is programmed using `tf.keras` it has the following

- 1 Define the model
- 2 Compile the model
- 3 Fit the model
- 4 Evaluate the model
- 5 Make predictions

Define the model


There are two ways to define the model

- 1 The sequential model
- 2 Functional API

The sequential model is what will be demonstrated

- A Sequential model is appropriate for a plain stack of layers where each layer has exactly one input tensor and one output tensor.
- The picture shows one "layer" but we could have multiple sequential layers in it, one thing follows the other.

```
model = tf.keras.Sequential(...)  
#array of "layers"
```



Source

Compile the Model

Compiling the model requires that you first select a loss (error/cost) function that you want to optimise, such as mean squared error or cross-entropy. You want parameters that give you the lowest error (controlling for overfitting/underfitting etc but that'll be covered next semester)

It also requires that you select an algorithm to perform the optimization procedure, typically stochastic gradient descent, or a modern variation, such as Adam.

The algorithm is typically an iterative method that will keep adjusting the weights until a minimum error is found.

```
model.compile(  
    optimizer=tf.optimizers.Adam(learning_rate=0.1),  
    loss='mean_absolute_error')
```


Fit the Model

Fitting the model means putting the data you have in and figuring out the "best" parameters for that data.

```
model.fit(X, y, epochs=500, verbose=0 )
```

where X is the array storing all the samples and their features and y is the vector storing the known response for each feature

Fitting is where the longest amount of time takes, in a neural network it can take weeks depending on the size of the dataset, the complexity of the model and the CPU/GPU power you've thrown at it. It also may require a lot of memory.

If we want to have a validation set to monitor at the end of each Epoch, we can add

```
model.fit(X, y, epochs=500, verbose=0, validation_split=0.2)
```

and you will get the `val_loss` and `val_accuracy` if you have specified these.

Warning

When using this for a validation set, 20% of the end of the data is just sliced away (no shuffling or randomisation!).

If your data is ordered information, you are going to get very poor results

If you have previously used `X_train, ... = train_test_split()`, then the data will be randomised so ok to use.

Evaluate the model

Often you will use a different dataset to evaluate the performance of the model rather than the one that you used to fit/train the model. This hopefully removes bias to the original set that was used to create the model.

Evaluating is much faster than fitting as the model is not changed but it is proportional to how much data you are using to evaluate

```
loss = model.evaluate(X, y, verbose=0)
```

And the whole point of everything! Being able to make predictions on previously unseen data.

In terms of code it is really easy. Let's say X_{new} is a vector of data (i.e. it has all the features for one sample then

```
yhat = model.predict(Xnew)
```

and it gives you the prediction. This is true whether we are using regression or classification (seeing if something is a picture of a cat/dog). You can also have X_{new} to be a matrix where each row is one sample.

Actually do an example

Move to workbook